

CSCI4180 (Fall 2021)

Assignment 1: Counting by MapReduce

Due on Oct 21, 2021, 23:59:59

Introduction

The goals of this assignment are to let students get familiar with Hadoop MapReduce and simple MapReduce development.

Part 1: Getting Started on Hadoop (20%)

In this part, you need to demonstrate the following:

1. Configure Hadoop in *fully distributed mode* on all your assigned virtual machines.
2. Run the provided WordCount program on your Hadoop platform.

Please show to the TAs that you achieve the above requirements.

We provide two sample datasets for you to test your program. One is the KJV Bible, and another one is the complete works of Shakespeare. Note that the sample datasets are not of large scale (each of them contains no more than 6MB), so you may not see the performance benefits when multiple mappers are used.

More instructions will be provided in the tutorials. Make sure that you strictly follow the instructions (e.g., the versions of Java and Hadoop, the configuration of the software); otherwise the TAs cannot troubleshoot the problems and you will risk losing marks!

In the following parts, we assume that the Hadoop cluster always runs in fully distributed mode.

Part 2: Word Length Count (20%)

Please extend the WordCount program to count the length of each word with the optimization technique *in-mapper combining*. Call the program *WordLengthCount.java*.

Output Format:

Each line contains a tuple of (length, count), separated by a space character. For example, “3 5” means that there are a total of five words of length three. The output is not necessarily printed in any increasing/decreasing order of length or count.

Sample Input:

```
The quick brown fox jumps  
over the lazy dog.
```

Sample Output:

```
3 3
5 3
4 3
```

Time Limit: Note that the program should be completed within a reasonable timeframe for a given dataset. The time limit is set by the TAs. Marks will be deducted if the program runs too long.

Note: The punctuations should also be counted toward the length of a word for your convenience. For example, in the sample input, the word “dog.” should be counted as a word with length 4 instead of 3.

Part 3: N -gram Initial Count (20%)

Using *WordLengthCount.java* as a starting point, extend it to count N -gram *initials* (call the program *NgramInitialCount.java*). An N -gram is a sequence of N consecutive words, where N is the input in the command-line argument. The *initials* of an N -gram are the first letters of the N words of the N -gram. In this problem, we consider a word as a sequence of English alphabet characters (e.g., dogs, cat, etc).

Output Format: Each line contains a tuple of (1st word’s initial, 2nd word’s initial, . . . , N -th word’s initial, count), where the adjacent items are separated by non-alphabet characters (e.g., space, tab, newline, punctuations, or numbers). You may assume that the datasets contain only printable ASCII characters, all in English. You can output the count of N -gram initials in any order. Words are case-sensitive. That is, “Apple apple” and “apple apple” should be counted and displayed separately. See the example below, where $N = 2$.

Command Format:

```
$ hadoop jar [.jar file] [class name] [input dir] [output dir] [N]
```

Sample Input:

```
‘The quick brown fox jumps
over the lazy dogs.
```

Sample Output:

```
T q 1
q b 1
b f 1
f j 1
j o 1
o t 1
t l 1
l d 1
```

Notes:

- We treat the punctuations “‘” and “.” as delimiters. We treat only “The” and “dogs” as words.
- For a word like “ca’t2e”, we treat it as three words “ca”, “t”, and “e”.
- The word in the end of a line and the word in the beginning of the next line can also form an N -gram. For example, “jumps over” is an N -gram for $N = 2$.
- Multiple non-alphabet characters are treated as a single delimiter. For example, “quick , , , , , brown” will form a 2-gram “quick brown”.
- We may miss some N -grams that appear in two different HDFS blocks, but it’s okay.
- You don’t need to implement in-mapper combining here, but you are strongly encouraged to think of any possible way to speed up your program.

Time Limit: Note that the program should be completed within a reasonable timeframe for a given

dataset. The time limit is set by the TAs. Marks will be deducted if the program runs too long.

Part 4: Computing N -gram Initials Relative Frequencies (20%)

Extend Part 3's program to compute the N -gram initials relative frequencies (call the program *NgramInitialRF.java*). Again, in the problem, you only need to consider the initial sequences whose initials are in all alphabets. For a given N , the relative frequency of an N -gram initial is defined as $\text{count}(\text{"X Y Z"})/\text{count}(\text{"X *"})$, where $\text{count}()$ is a function that counts the occurrence of an input, $X Y Z$ are the initials of an N -gram (for $N = 3$ here), and $*$ stands for any alphabet initial(s).

Output Format: Each line contains a tuple of (1st word's initial, 2nd word's initial, \dots , N -th word's initial, frequency). You only need to output the sequences whose relative frequencies are at least θ , where θ is a configurable parameter specified in the command-line argument and $0 \leq \theta \leq 1$. See the following example for $N = 2$ and $\theta = 0.6$.

Command Format:

```
$ hadoop jar [.jar file] [class name] [input dir] [output dir] [N] [ $\theta$ ]
```

Sample Input:

```
the quick brown fox jumps over the lazy dogs.
the Quick Brown fox jumps Over the Lazy dogs.
the Quick Brown fox Jumps Over the Lazy 'dogs'.
```

Sample Output:

```
f j 0.666666667
q b 1
b f 1
o t 1
l d 1
d t 1
Q B 1
B f 1
O t 1
L d 1
J O 1
```

Time Limit: Note that the program should be completed within a reasonable timeframe for a given dataset. The time limit is set by the TAs. Marks will be deducted if the program runs too long.

Notes:

- For the 2-gram initial "f j", since $\text{count}(\text{"f j"}) = 2$ and $\text{count}(\text{"f *"}) = 3$, the relative frequency is $2/3 > \theta$.
- For the N -gram initial "l d", we can see that $\text{count}(\text{"l d"}) = 1$ and $\text{count}(\text{"l *"}) = 1$, so the relative frequency is 1.
- For the precision of relative frequency, both float and double data types are okay.

Part 5: Configure Hadoop on Amazon EC2 (20%)

In this part, you need to demonstrate the following:

1. Configure Hadoop in fully distributed mode on Amazon EC2, using 3 compute instances.

2. Run the provided WordCount program on EC2.

Please show to the TAs that you achieve the above requirements.

Bonus (5%)

The top 3 groups whose Part 4's programs have the smallest running time will receive the bonus marks. You may consider to optimize your programs or configure some parameters in Hadoop to make the programs perform better. If more than 3 groups have the best performance, we will still give out the 5% bonus to each group. Note that the program must return the correct answer in order to be considered for the bonus mark.

Submission Guidelines

You *must* at least submit the following files, though you may submit additional files that are needed:

- *WordLengthCount.java*
- *NgramInitialCount.java*
- *NgramInitialRF.java*
- Declaration form for group projects:
([http://www.cuhk.edu.hk/policy/academichonesty/Eng_hm_files_\(2013-14\)/p10.htm](http://www.cuhk.edu.hk/policy/academichonesty/Eng_hm_files_(2013-14)/p10.htm))

Your programs must be implemented in Java. Any program that is not compilable will receive zero marks!!

The testing platform during demo is provided by the TAs. If you need to change the configuration of Hadoop, make sure that you do it dynamically inside the code. You are not allowed to modify any configuration files during demo.

Please refer to the course website for the submission instructions.

Demo will be arranged on the next day.

Have fun! :)